

Session 4



CTF Track: Reverse Engineering

Reminder

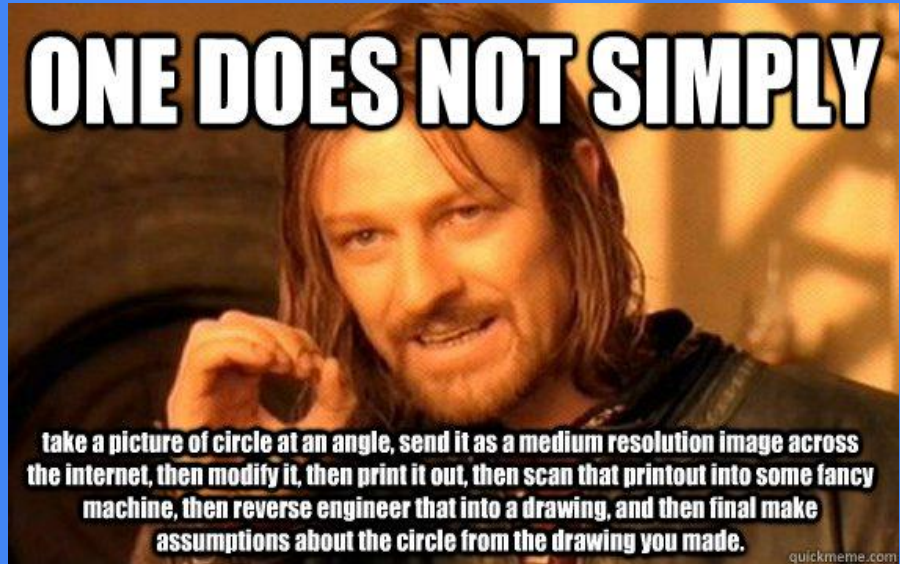
As announced last time, we will participating in TU CTF on Nov 26th from 10am to 3pm.

Contents

- What is reverse engineering?
- Basic skills
- Practice problems

What is reverse engineering?

Oops.



What is reverse engineering?

(in general)

Wikipedia has a pretty abstract definition:

Reverse engineering, also called back engineering, is the processes of extracting knowledge or design information from a product and reproducing it or reproducing anything based on the extracted information.

What is reverse engineering?

(for software)

Most software is licensed to you under the premise that you do not have the rights to reverse engineer it. It basically means you are not allowed to run it inside a debugger to discover how it works and to modify it.

M. No Reverse Engineering. You may not, and you agree not to or enable others to, copy (except as expressly permitted by this License or by the Usage Rules if they are applicable to you), decompile, reverse engineer, disassemble, attempt to derive the source code of, decrypt, modify, or create derivative works of the Apple Software or any services provided by the Apple Software or any part thereof (except as and only to the extent any foregoing restriction is prohibited by applicable law or by licensing terms governing use of Open-Sourced Components that may be included with the Apple Software).

**An excerpt from the License Agreement of
some Apple software**

What is reverse engineering?

(for security researchers)

Security researchers ignore those rules.

They reverse engineer software as they see fit, usually to discover and report security vulnerabilities.

bugs.chromium.org/project-zero/issues/list?can=1&redir=1

Issues - project-zero - Project Zero - Monorail

Project: project-zero Issues People Development process History

New issue Search New issues for Search Advanced search Search tips

1 - 100 of 1199 Next List Grid

ID	Type	Status	Priority	Milestone	Owner	Summary + Labels
1	Invalid	---	---	---	cevas@google.com	This is a test
9	Fixed	---	---	---	cevas@google.com	Safari sandbox logic error enables reading of arbitrary files
10	Fixed	---	---	---	cevas@google.com	Safari sandbox IPC memory corruption with WebEvent:Wheel
11	Fixed	---	---	---	cevas@google.com	Safari sandbox IPC memory corruption with WebEvent:Char
12	Fixed	---	---	---	cevas@google.com	launchd heap corruption due to integer overflow in launch_data_unpack
13	Fixed	---	---	---	cevas@google.com	launchd heap corruption due to incorrect rounding in launch_data_unpack
14	Fixed	---	---	---	cevas@google.com	launchd heap overflow in log_forward
15	Fixed	---	---	---	cevas@google.com	Lack of bounds checking in notified CCProjectZeroMembers
16	Fixed	---	---	---	cevas@google.com	launchd heap corruption due to unchecked stropy in init_session MIG ipc
17	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to lack of bounds checking in IOAccel2DContext2-bit
18	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel memory disclosure due to lack of bounds checking in AGPMClient:getPstatesOccupancy
19	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to unchecked pointer parameter in IOAccelCLContext:unmap_user_memory
20	Fixed	---	---	---	cevas@google.com	OS X IOKit Multiple exploitable kernel NULL dereferences (x4)
21	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel memory disclosure due to lack of bounds checking in IOUSBControllerUserClient:ReadRegister
22	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to incorrect bounds checking in Intel GPU driver (x2)
23	Fixed	---	---	---	cevas@google.com	OS X KASLR defeat using sgdt
24	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to NULL pointer dereference in IOTThunderboltFamily
28	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to lack of bounds checking in GPU command buffers
29	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to off-by-one error in IOAccelGLContext:processSidebandToken
30	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel multiple exploitable memory safety issues in token parsing in IOAccelVideoContextMedia (x5)
31	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to NULL pointer dereference in IOAccelContext2:clientMemoryForType
32	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to lack of bounds checking in IOAccelVideoContextMain:process_token_ColorSpaceConversion
33	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to lack of bounds checking in IOAccelDisplayPipeTransaction2:set_plane_gamma_table
34	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to multiple bounds checking issues in IOAccelGLContext token parsing (x3)
35	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to controlled mem_free size in IOSharedDataQueue
36	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to lack of bounds checking in AppleMultitouchIODataQueue
37	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to bad free in IOBluetoothFamily
38	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to integer overflow in IOBluetoothDataQueue (root only)
39	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to integer overflow in IODataQueue:enqueue
40	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to heap overflow in IOHKeyboardMapper:parseKeyMapping
41	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel code execution due to NULL pointer dereference in IOHKeyboardMapper:stickyKeysFree
42	Fixed	---	---	---	cevas@google.com	OS X IOKit kernel memory disclosure due to lack of bounds checking in IOHKeyboardMapper:modifierSwapFilterKey
43	Fixed	---	---	---	cevas@google.com	Flash leak of uninitialized data whilst rendering JPEGs
44	Fixed	---	---	---	cevas@google.com	Flash leak of uninitialized data whilst rendering a 2-component JPEG
45	Fixed	---	---	---	cevas@google.com	Flash leak of uninitialized memory when rendering valid(?) 1bpp image
46	Fixed	---	---	---	cevas@google.com	Flash heap buffer overflow calling copyPixelsToByteArray() on a large ByteArray CCProjectZeroMembers
47	Fixed	---	---	---	cevas@google.com	Flash leak of uninitialized data when image zlib stream ends prematurely CCProjectZeroMembers
48	Fixed	---	---	---	cevas@google.com	Flash leak of uninitialized data when JPEG image alpha channel zlib stream ends prematurely CCProjectZeroMembers
71	Fixed	---	---	---	cevas@google.com	Flash out-of-bounds read in uploadCompressedTextureFromByteArray() CCProjectZeroMembers

Google's "Project Zero" is basically all about reverse engineering to find vulns

What is reverse engineering?

(for CTFs)

In CTFs, reverse engineering questions usually involve:

1. Being provided an unknown binary
2. Find out some way to coerce it into doing something

Typical CTF Problem

Here is a binary that performs some funky operations on a password you specify, and then after these funky operations it may give you the flag.

Basic Skills


In beginner CTFs, sometimes the source code to that unknown executable will be provided. So you will read the source code to figure out how things work.

Also in beginner CTFs, sometimes the flag as a string is inside the executable. This is when `strings` is helpful.

**I HEARD YOU
LIKE
STRINGS.**

STRING THEORY SUMMARIZED:

I JUST HAD AN AWESOME IDEA.
SUPPOSE ALL MATTER AND ENERGY
IS MADE OF TINY, VIBRATING "STRINGS."



Find the printable strings in a binary.

Find the printable strings in a binary.

What is `strings(1)`?

strings

Its man page is less than one page (on a Mac; two pages on Linux). Just read it.

STRINGS(1)

STRINGS(1)

NAME

strings – find the printable strings in a object, or other binary, file

SYNOPSIS

strings [-] [-a] [-o] [-t *format*] [-number] [-n *number*] [--] [file ...]

DESCRIPTION

Strings looks for ASCII strings in a binary file or standard input. *Strings* is useful for identifying random object files and many other things. A string is any sequence of 4 (the default) or more printing characters [ending at, but not including, any other character or EOF]. Unless the - flag is given, *strings* looks in all sections of the object files except the (__TEXT,__text) section. If no files are specified standard input is read.

The file arguments may be of the form *libx.a(foo.o)*, to request information about only that object file and not the entire library. (Typically this argument must be quoted, "*libx.a(foo.o)*", to get it past the shell.)

The options to *strings*(1) are:

- a This option causes *strings* to look for strings in all sections of the object file (including the (__TEXT,__text) section).
- This option causes *strings* to look for strings in all bytes of the files (the default for non-object files).
- This option causes *strings* to treat all the following arguments as files.
- o Preceded each string by its offset in the file (in decimal).
- t *format* Write each string preceded by its byte offset from the start of the file. The format shall be dependent on the single character used as the format option-argument:
 - d The offset shall be written in decimal.
 - o The offset shall be written in octal.
 - x The offset shall be written in hexadecimal.
- number The decimal *number* is used as the minimum string length rather than the default of 4.
- n *number* Specify the minimum string length, where the number argument is a positive decimal integer. The default shall be 4.
- arch *arch_type* Specifies the architecture, *arch_type*, of the file for *strings*(1) to operate on when the file is a universal file. (See *arch*(3) for the currently known *arch_types*.) The *arch_type* can be "all" to operate on all architectures in the file.

SEE ALSO

od(1)

BUGS

The algorithm for identifying strings is extremely primitive.

Simple Usage

```
$ strings -n 8 funny < /path/to/your/file
```

The strings utility is by default prints any 4 consecutive printable characters in a binary file.

But this frequently results in false positives because typical x86-64 machine code can contain many sequences of 4 consecutive printable characters.

We recommend passing `-n 8` (or maybe `-n 6` if you don't consider yourself lucky).

Exploration Time!

Try to run the following and see what you can find:

- `strings -n 8 $(which cat) # find printable strings in the cat(1) utility`
- `strings -n 8 $(which ls) # find printable strings in the ls(1) utility`
- `strings -n 8 $(which cp) # find printable strings in the cp(1) utility`

What do you see? Why do you see those strings? What do you think these strings are used for?

Discuss with people around you (5 mins).

The Hunt for Paul Eggert

Challenge: for all executables in your PATH, run `strings -n 8` on them. How many of them contains the string “Eggert”? Which?

The Hunt for Paul Eggert: Example Solution

```
for p in $(echo $PATH | tr : $'\n'); do  
    find $p -maxdepth 1 -type f -exec sh -c 'strings -n 8 {} | fgrep  
-q Paul\ Eggert' \; -print  
  
done 2> /dev/null
```

On a SEASnet machine with typical PATH configuration, I find 20 of them.

Now manually run strings to find out why they have “Paul Eggert.”

Alternatives to strings

Although strings is simple and popular, many more powerful utilities can perform similar functionalities of finding and extracting printed ASCII characters, including old good grep and its extended cousin egrep.

grep(1)/egrep(1) allows you to only print matched fragments within a file, without dumping all the binary and unprintable characters.

grep(1)/egrep(1) also allows recursive searching without using tools like find(1) or xargs(1).

The Hunt for Paul Eggert, revisited

Challenge: for all files (recursively) in /usr/bin, which of them contains a string that has an uppercase P, followed by three alphabetical characters, a space, an uppercase “E”, followed by four alphabetical characters, followed by a lowercase “t”?

What are the matched strings in question?

The Hunt for Paul Eggert, revisited: Example Solution

To find which files:

```
egrep -rl 'P[[:alpha:]]{3} E[[:alpha:]]{4}t' /usr/bin 2>/dev/null
```

To find what the matched strings are:

```
egrep -aro 'P[[:alpha:]]{3} E[[:alpha:]]{4}t' /usr/bin 2>/dev/null
```

Packers (Executable compressor)

Even though a CTF reverse engineering problem could be more complicated than just running strings, sometimes it can reveal important clues.

Take for example flag from pwnable.kr. When you run strings, one of the lines is:

```
$Info: This file is packed with the UPX executable packer  
http://upx.sf.net $  
$Id: UPX 3.08 Copyright (C) 1996-2011 the UPX Team. All  
Rights Reserved. $
```

This tells you the file is packed with UPX and you can use `upx -d` to unpack it.

**Is every CTF rev
eng problem as
simple as
running strings?
I sure hope so.**

(It's not.)

What if source code is provided?

Since many of you haven't taken CS33 yet, let's first tackle the kind of problems where source code is given.

Take collision from pwnable.kr for example.

Walkthrough: pwnable.kr collision

```
#include <stdio.h>
#include <string.h>
unsigned long hashcode = 0x21DD09EC;
unsigned long check_password(const char* p){
    int* ip = (int*)p;
    int i;
    int res=0;
    for(i=0; i<5; i++){
        res += ip[i];
    }
    return res;
}

int main(int argc, char* argv[]){
    if(argc<2){
        printf("usage : %s [passcode]\n", argv[0]);
        return 0;
    }
    if(strlen(argv[1]) != 20){
        printf("passcode length should be 20 bytes\n");
        return 0;
    }

    if(hashcode == check_password( argv[1] )){
        system("/bin/cat flag");
        return 0;
    }
    else
        printf("wrong passcode.\n");
    return 0;
}
```

Understanding the main function

- It first makes sure one argument is provided.
- Then it makes sure it is exactly 20 bytes.
- Then it makes sure after doing some funky operations, the result is `0x21DD09EC`.
- Then it gives you the flag!

Understanding the check_password function

```
unsigned long check_password(const char* p){  
    int* ip = (int*)p;  
    int i;  
    int res=0;  
    for(i=0; i<5; i++){  
        res += ip[i];  
    }  
    return res;  
}
```

- It first casts a pointer to characters to a pointer to int. Essentially reinterpret_cast in C++.
- And it adds together the five integers that are contained in those 20 bytes.

Understanding the check_password function

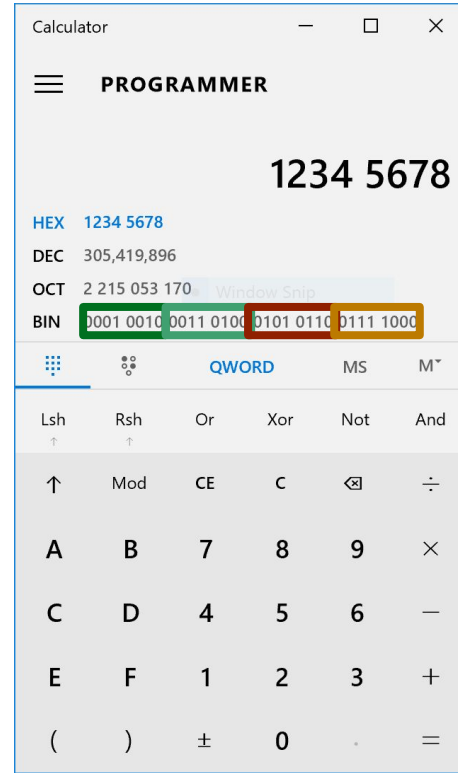
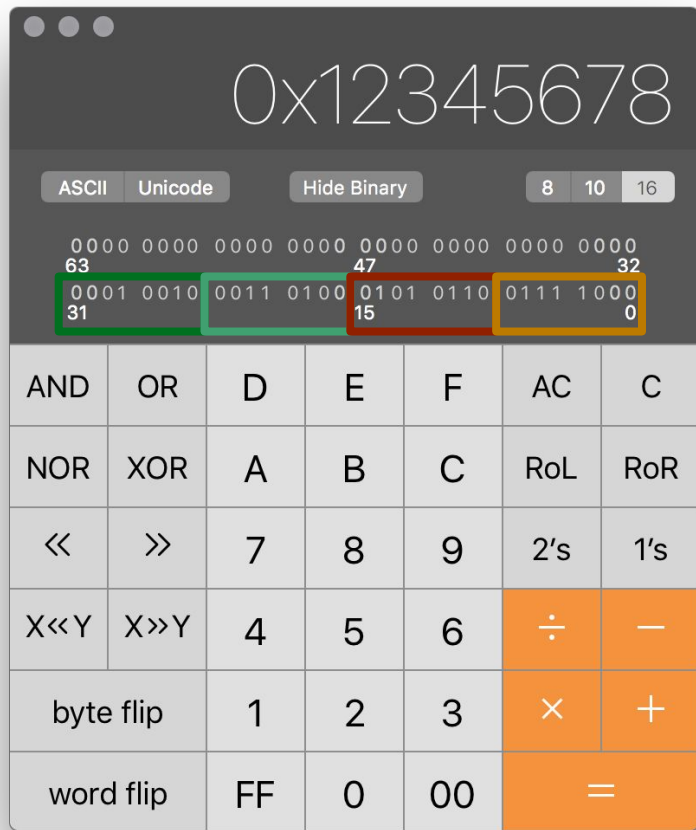


- Each small box above is a character you need to provide.
- Each colored region is a reinterpreted integer.

Understanding the check_password function



- Looking at one integer.
- Each colored region is one byte now.
- Little endian: least significant byte on the left, most significant byte on the right.
- Remember to use your calculator to help you!



Use your calculator! Remember however that it uses big-endian so the order is reversed.

Advanced Reverse Engineering: x86-64 and disassemblers

In more advanced CTF problems, you will not be given the source code.

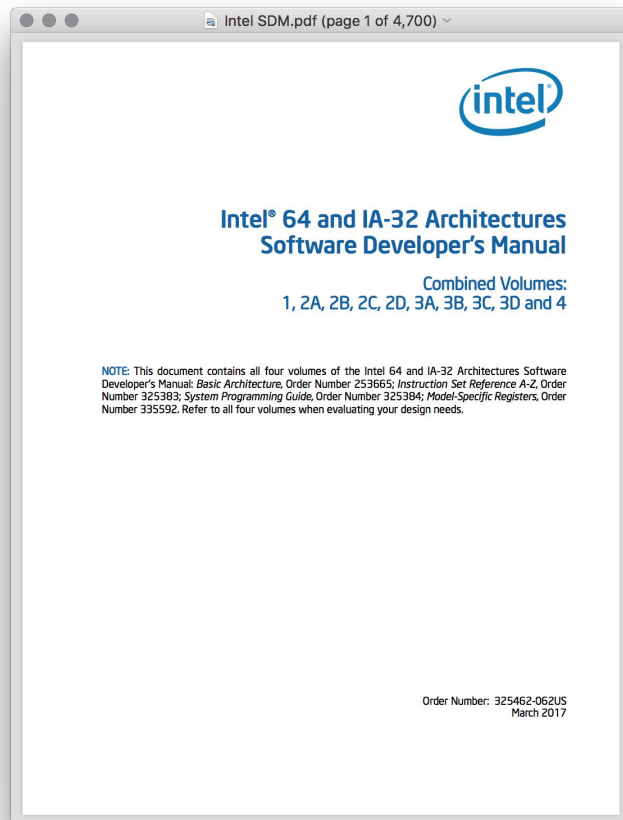
You then have to use disassemblers and/or debugger such as GDB.

You will also need to use GDB.

The x86 and x86-64 architecture

Very few people on few know every detail of x86-64 systems. The full documentation is a staggering 4,700 pages.

(It was my bedtime reading when I was taking CS33).

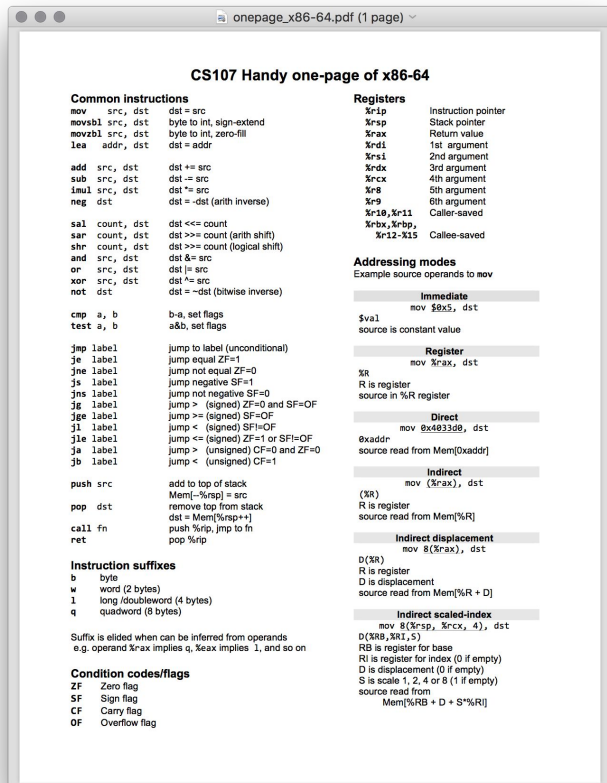


The x86 and x86-64 architecture

Basic concepts:

- What is RAM?
 - What's the stack?
 - How do we address memory?
- What are registers?
- What is code?
- How do C control structures map to assembler?
 - If/else, function calls, loops, switches

Use this one-page summary of x86-64 from Stanford as a reference!



tinyurl.com/CTEproblems4

tinyurl.com/CTEfeedback4